

# Introduction to Deep Learning (I2DL)

Solution to exercise 3 and exam tips

# Today

- Solution of Exercise 3
  - Facial Keypoints
  - Recurrent Neural Networks
  - Segmentation
- Practical Pytorch Stuff
- Exam Topics and Mock Exam
- Organizationional Stuff

# Solutions to EX3

# Spatial Batchnormalization

- Forward

```
424 #####  
425  
426 # Computation in one sweep by rearranging the dims to fit into  
427 # the batchnorm_forward framework  
428 x_swapped = np.transpose(x, (0, 2, 3, 1))  
429 x_swapped_reshaped = np.reshape(x_swapped, (-1, x_swapped.shape[-1]))  
430  
431 out_temp, cache = batchnorm_forward(  
432     x_swapped_reshaped, gamma, beta, bn_param)  
433 out = np.transpose(np.reshape(out_temp, x_swapped.shape), (0, 3, 1, 2))  
434  
435 #####
```

- Backward

```
463 #####  
464  
465 dout_swapped = np.transpose(dout, (0, 2, 3, 1))  
466 dout_swapped_reshaped = np.reshape(  
467     dout_swapped, (-1, dout_swapped.shape[-1]))  
468  
469 dx_sr, dgamma, dbeta = batchnorm_backward(dout_swapped_reshaped, cache)  
470  
471 dx = np.transpose(np.reshape(dx_sr, dout_swapped.shape), (0, 3, 1, 2))  
472  
473 #####
```

# A word of warning

- Be aware of the differences of regular and spatial Batchnormalization and Dropout
- Pytorch will usually not throw an error as it is a valid possibility

E.g., for 2d tasks use:

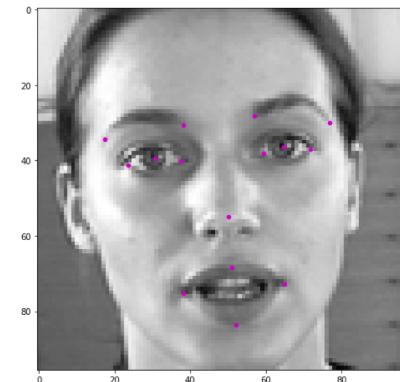
- Dropout2d
- BatchNorm2d

# Facial Keypoints

- Facial Keypoint extraction is a Regression task



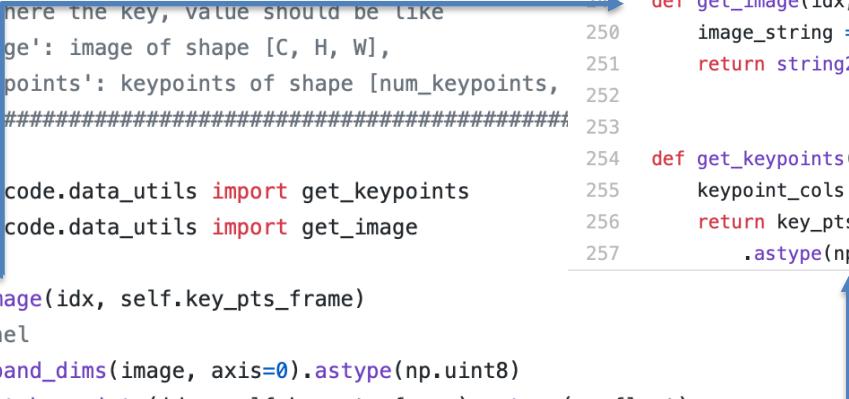
Input  
Image



Coordinates  
Facial  
Keypoints

# Facial Keypoints Dataloader

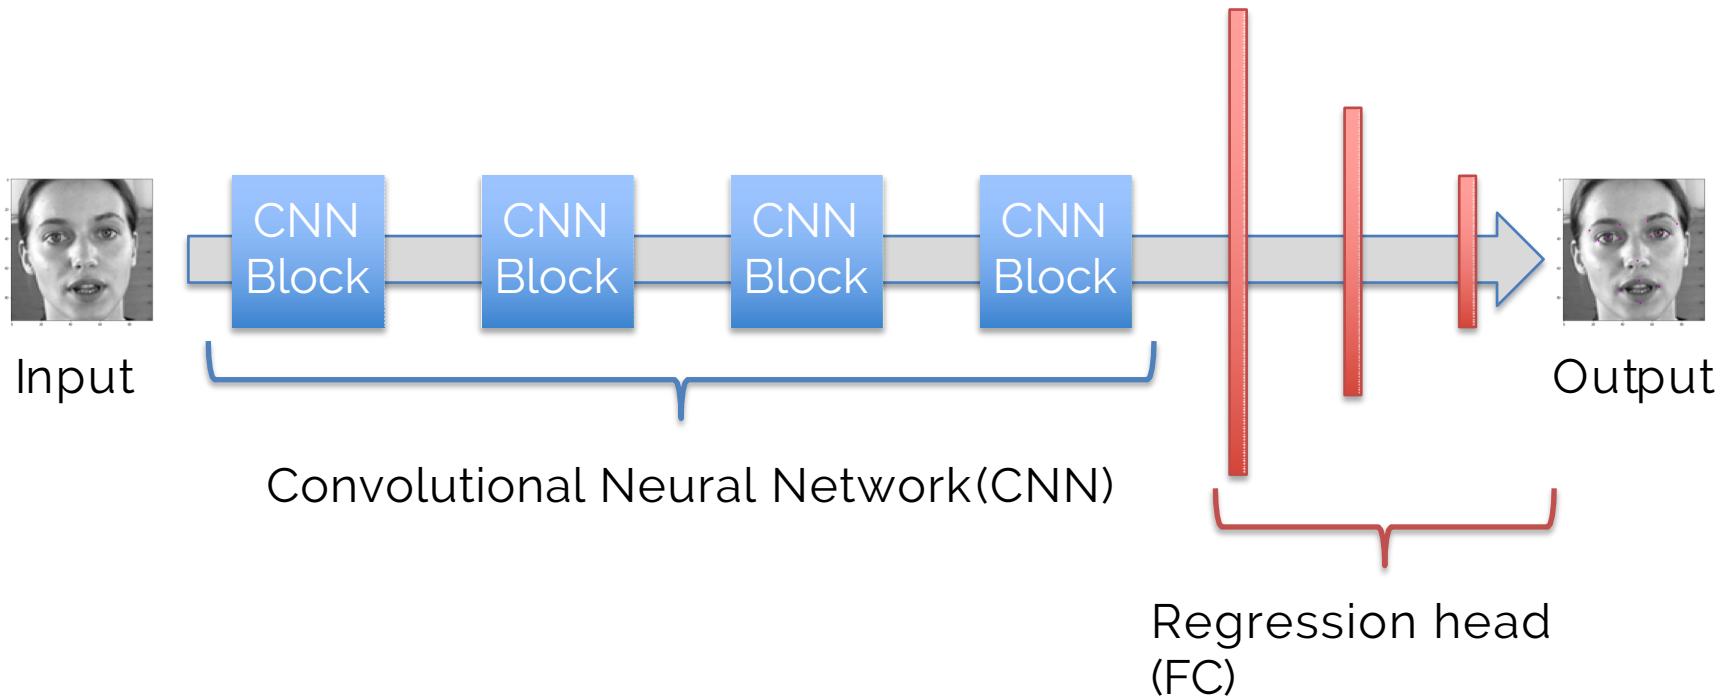
```
--  
35     def __getitem__(self, idx):  
36         sample = {'image': None, 'keypoints': None}  
37         #####  
38         # TODO:  
39         # Return the idx sample in the Dataset. A sample should be a  
40         # dictionary where the key, value should be like 248  
41         #       {'image': image of shape [C, H, W], 250  
42         #       'keypoints': keypoints of shape [num_keypoints, 251  
43         #       ##### 252  
44  
45         from exercise_code.data_utils import get_keypoints  
46         from exercise_code.data_utils import get_image  
47  
48         image = get_image(idx, self.key_pts_frame)  
49         # expand channel  
50         image = np.expand_dims(image, axis=0).astype(np.uint8)  
51         keypoints = get_keypoints(idx, self.key_pts_frame).astype(np.float) 253  
52  
53         sample = {'image': image, 'keypoints': keypoints}  
54  
55         #####  
  
248     #  
249     #  
250     def get_image(idx, key_pts_frame):  
251         image_string = key_pts_frame.loc[idx]['Image']  
252         return string2image(image_string)  
253  
254     def get_keypoints(idx, key_pts_frame):  
255         keypoint_cols = list(key_pts_frame.columns)[-1]  
256         return key_pts_frame.iloc[idx][keypoint_cols].values.reshape((15, 2))\br/>257         .astype(np.float32)
```



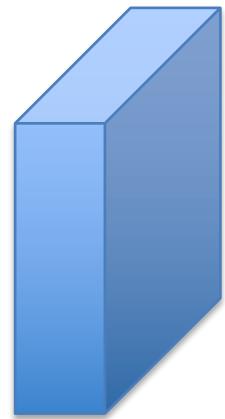
# Transforms

```
5  class Normalize(object):
6      """
7      Normalizes keypoints.
8      """
9      def __call__(self, sample):
10
11         image, key_pts = sample['image'], sample['keypoints']
12
13         #####
14         # TODO: Implement the Normalize function, where we normalize #
15         # the image from [0, 255] to [0,1] and keypoints from [0, 96]#
16         # to [-1, 1]                                              #
17         #####
18
19         key_pts = (key_pts - 48.) / 48.
20         image = image / 255.0 *2 - 1
21
22         #####
23         # End of your code                                     #
24         #####
25
26         return {'image': image, 'keypoints': key_pts}
```

# Network Architecture



# CNN Block



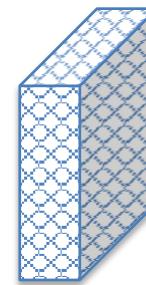
Conv  
Layer



Max  
Pooling



ReLU



(2dDropout)

# Input / Output size conv layer

$$W_O = (W_I - F + 2P) / S + 1$$

Output Size  
Input Size  
Filter Size  
Zero Padding  
Stride

Stride = 1

Padding = 1

You should  
know!

# Network Architecture

$$W_O = (W_I - F + 2P) / S + 1$$

↑      ↑      ↑      ↑      ↑  
Output Size    Input Size    Filter Size    Zero Padding    Stride

(1x 96x 96)    (32x 48x 48)    (64x 24x 24)    (128x 12x 12)    (256x 6x 6)



Input

CNN Block

CNN Block

CNN Block

CNN Block

Stride =1  
Padding =1

Convolutional Neural Network(CNN)

# Network Code

```
24
25     def conv_sandwich(inp, out, kernel_size, stride, pad):
26         return nn.Sequential(
27             nn.Conv2d(inp, out, kernel_size, stride, pad),
28             nn.MaxPool2d(2, 2),
29             nn.ReLU()
30         )
31
32     layers = []
33     layers.append(conv_sandwich(1, 32, kernel_size=3, stride=1, pad=1))
34     layers.append(conv_sandwich(32, 64, kernel_size=3, stride=1, pad=1))
35     layers.append(conv_sandwich(64, 128, kernel_size=3, stride=1, pad=1))
36     layers.append(conv_sandwich(128, 256, kernel_size=3, stride=1, pad=1))
37     self.convs = nn.Sequential(*layers)
38
39     self.fc1 = nn.Sequential(nn.Linear(256 * 6 * 6, 256), nn.ReLU())
40     self.fc2 = nn.Sequential(nn.Linear(256, 30), nn.Tanh())
41
```

# Keypoints as Heatmaps



(a)



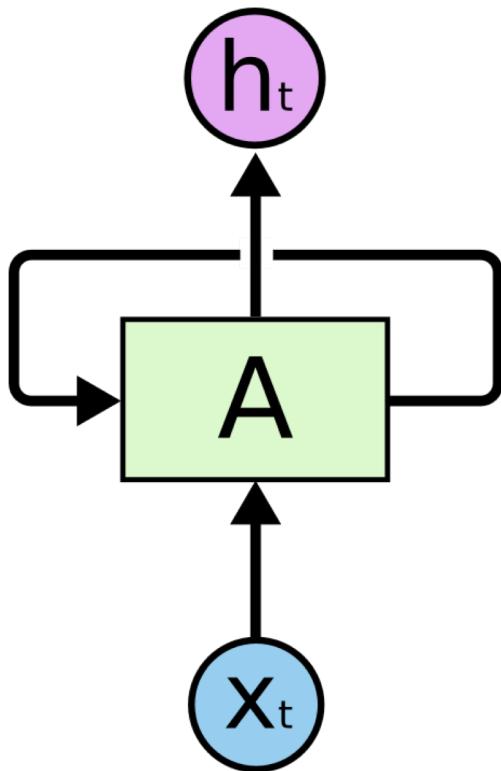
(b)



(c)

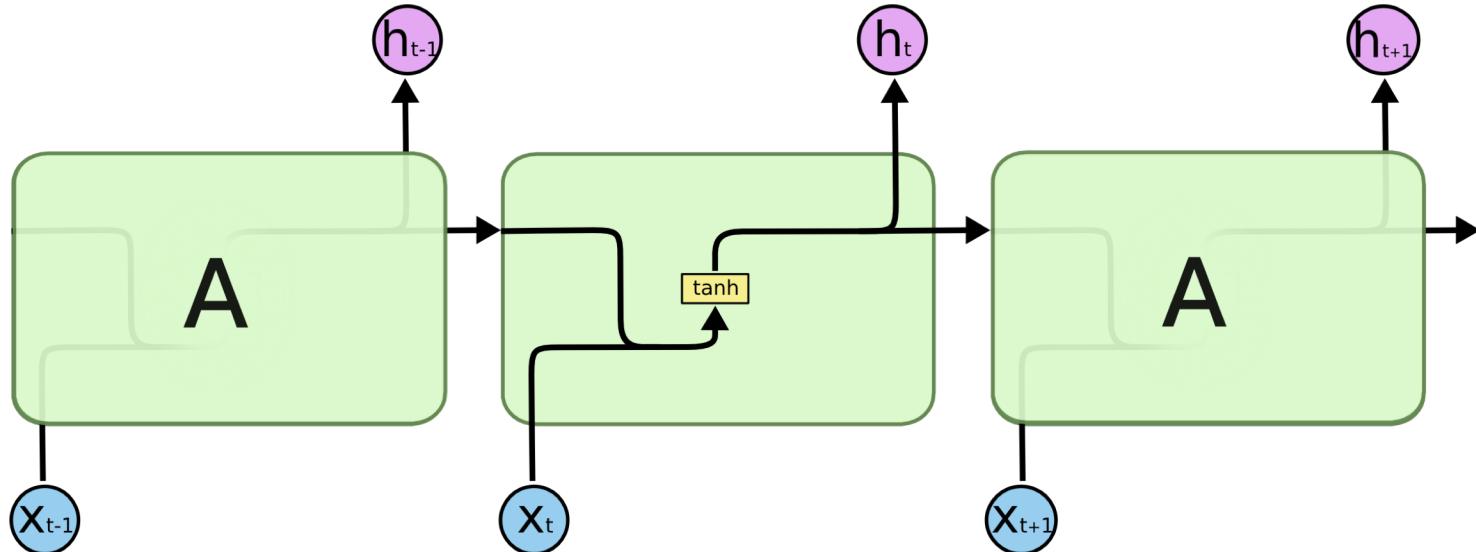
Loss in this situation: Softmax

# Recurrent Neural Nets



# Simple RNN

$$h_t = \tanh (W \cdot h_{t-1} + V \cdot x_t + b)$$



# Simple RNN

```
21     self.hidden_size = hidden_size
22     self.input_size = input_size
23
24     self.W_hh = nn.Linear(self.hidden_size, self.hidden_size, bias=True)
25     self.W_xh = nn.Linear(self.input_size, self.hidden_size, bias=True)
26     if activation == "tanh":
27         self.activation = nn.Tanh()
28     elif activation == "relu":
29         self.activation = nn.ReLU()
30
31     else:
32         raise ValueError(
33             "Unrecognized activation. Allowed activations: tanh or relu")
```

# Simple RNN

- Forward

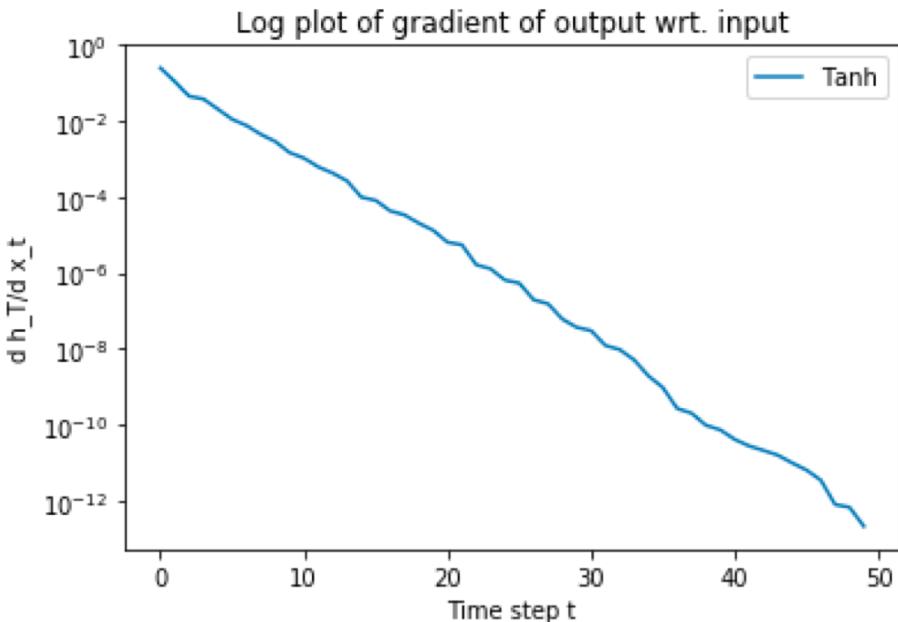
```
55     if h is None:  
56         h = torch.zeros((1, x.size(1), self.hidden_size)).float()  
57  
58     h_seq = []  
59     for t in range(x.size(0)):  
60         # update the hidden state  
61  
62         h = self.activation(self.W_hh(h) + self.W_xh(x[t]))  
63         h_seq.append(h)  
64     h_seq = torch.cat(h_seq, 0)
```

- Own RNN vs Pytorch:

- Time Pytorch RNN 10000 runs: 2.796s
- Time I2DL RNN 10000 run: 6.434s

# Simple RNN

- Vanishing gradient



$$h_t = \tanh(W \cdot h_{t-1} + V \cdot x_t + b)$$

$$\sigma(\cdot) = \tanh(\cdot) \quad z_t = W \cdot h_{t-1} + V \cdot x_t$$

$$\frac{\partial h_t}{\partial x_\tau} = \frac{\partial \sigma(z_t)}{\partial z_t} \left( W \cdot \frac{\partial h_{t-1}}{\partial x_\tau} + V \cdot \delta_{t\tau} \right)$$

# Simple RNN

$$\boxed{\begin{aligned} h_t &= \tanh(W \cdot h_{t-1} + V \cdot x_t + b) \\ \sigma(\cdot) &= \tanh(\cdot) \quad z_t = W \cdot h_{t-1} + V \cdot x_t \end{aligned}}$$

- Calculate  $\frac{\partial h_3}{\partial x_0}$  with  $h_0 = 0$

$$\frac{\partial h_t}{\partial x_\tau} = \frac{\partial \sigma(z_t)}{\partial z_t} \left( W \cdot \frac{\partial h_{t-1}}{\partial x_\tau} + V \cdot \delta_{t\tau} \right)$$

$$\frac{\partial h_3}{\partial x_1} = \frac{\partial \sigma(z_3)}{\partial z_3} \left( W \cdot \frac{\partial \sigma(z_2)}{\partial z_2} \left( W \cdot \frac{\partial \sigma(z_1)}{\partial z_1} \cdot V \right) \right)$$

-> Gradient depends on the norm of  $W$

# LSTM

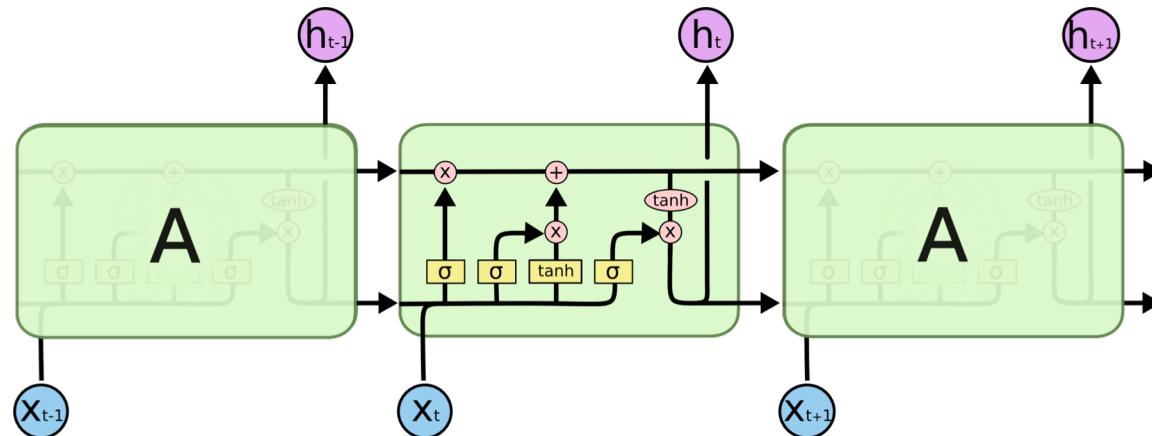
$$f_t = \sigma_g (W_f \cdot h_{t-1} + V_f \cdot x_t + b_f)$$

$$i_t = \sigma_g (W_i \cdot h_{t-1} + V_i \cdot x_t + b_i)$$

$$o_t = \sigma_g (W_o \cdot h_{t-1} + V_o \cdot x_t + b_o)$$

$$c_t = f_t \circ c_{t-1} + i_t \circ \tanh (W_c \cdot h_{t-1} + V_c \cdot x_t + b_c)$$

$$h_t = o_t \circ \tanh (c_t)$$



# LSTM

- Initialisation

```
80 #####  
81  
82     self.hidden_size = hidden_size  
83     self.input_size = input_size  
84  
85     self.W_hh = nn.Linear(self.hidden_size, 4*self.hidden_size, bias=True)  
86     self.W_xh = nn.Linear(self.input_size, 4*self.hidden_size, bias=True)  
87 #####  
88 #####
```

# LSTM

- Forward

```
110     if h is None:
111         h = torch.zeros((1, x.size(1), self.hidden_size)).float()
112     if c is None:
113         c = torch.zeros((1, x.size(1), self.hidden_size)).float()
114     h_seq = []
115     for t in range(x.size(0)):
116         # update the hidden state
117         update_h = self.W_hh(h)
118         update_x = self.W_xh(x[t].unsqueeze(0))
119
120         gates = (update_h[:, :, : 3 * self.hidden_size] +
121                  update_x[:, :, : 3 * self.hidden_size]).sigmoid()
122         fg = gates[:, :, :self.hidden_size]
123         ig = gates[:, :, self.hidden_size:2 * self.hidden_size]
124         og = gates[:, :, 2 * self.hidden_size:]
125         update = (update_h[:, :, 3 * self.hidden_size:] +
126                   update_x[:, :, 3 * self.hidden_size:]).tanh()
127         c = c * fg + update * ig
128         h = og * c.tanh()
129
130         h_seq.append(h)
131     h_seq = torch.cat(h_seq, 0)
```

Own LSTM vs Pytorch:

Time Pytorch LSTM 10000 runs: **6.669s**  
Time I2DL LSTM 10000 run: **16.756s**

# LSTM

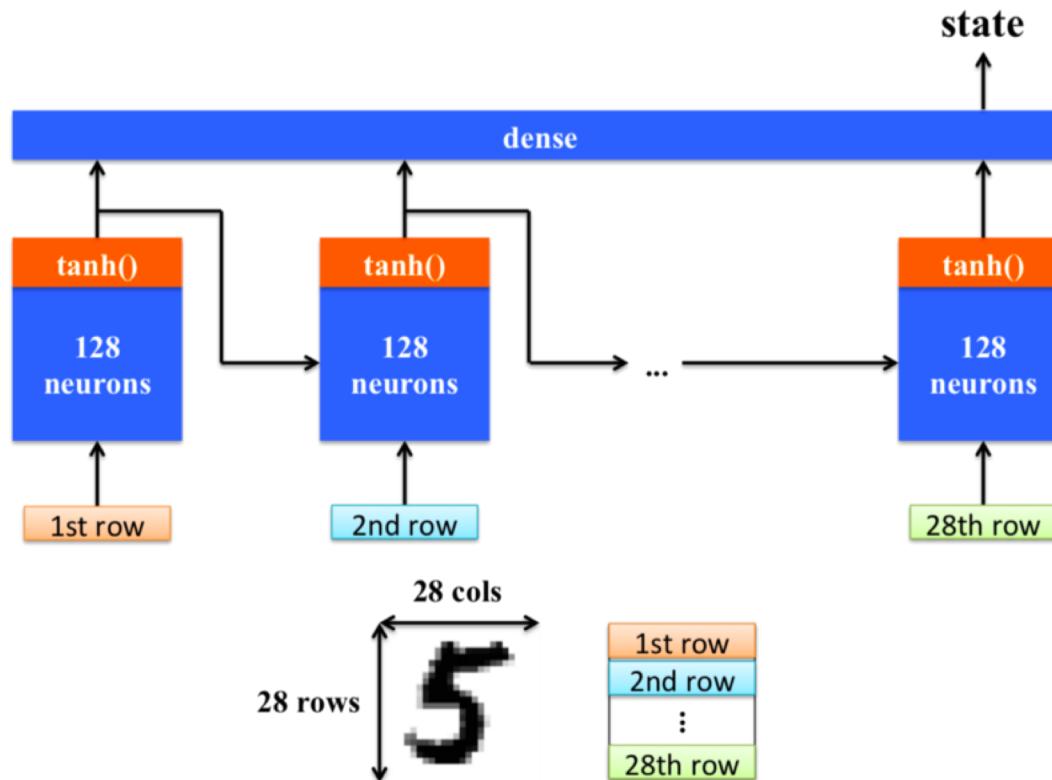
- LSTM "resolves" vanishing gradient problem

```
hidden_size=1  
input_size= 1  
  
time_steps=50  
rnn=RNN(input_size, hidden_size)  
x = torch.randn(time_steps, 1, input_size)  
x.requires_grad=True  
,h=rnn(x)  
h.requires_grad  
h.sum().backward()  
grad_rnn=X.grad.view(-1)  
  
lstm=LSTM(input_size, hidden_size)  
x = torch.randn(time_steps, 1, input_size)  
x.requires_grad=True  
,(h, c)=lstm(x)  
h.sum().backward()  
grad_lstm=x.grad.view(-1)
```



```
plt.semilogy(np.flip(abs(grad_lstm.detach().cpu().numpy()), 1), label="LSTM")  
plt.semilogy(np.flip(abs(grad_rnn.detach().cpu().numpy()), 1), label="RNN")  
plt.legend()  
plt.xlabel("Time step t")  
plt.ylabel("d h_T/d x_t")  
plt.title("Log plot of gradient of output wrt. input")  
plt.show()
```

# MNIST RNN Classifier

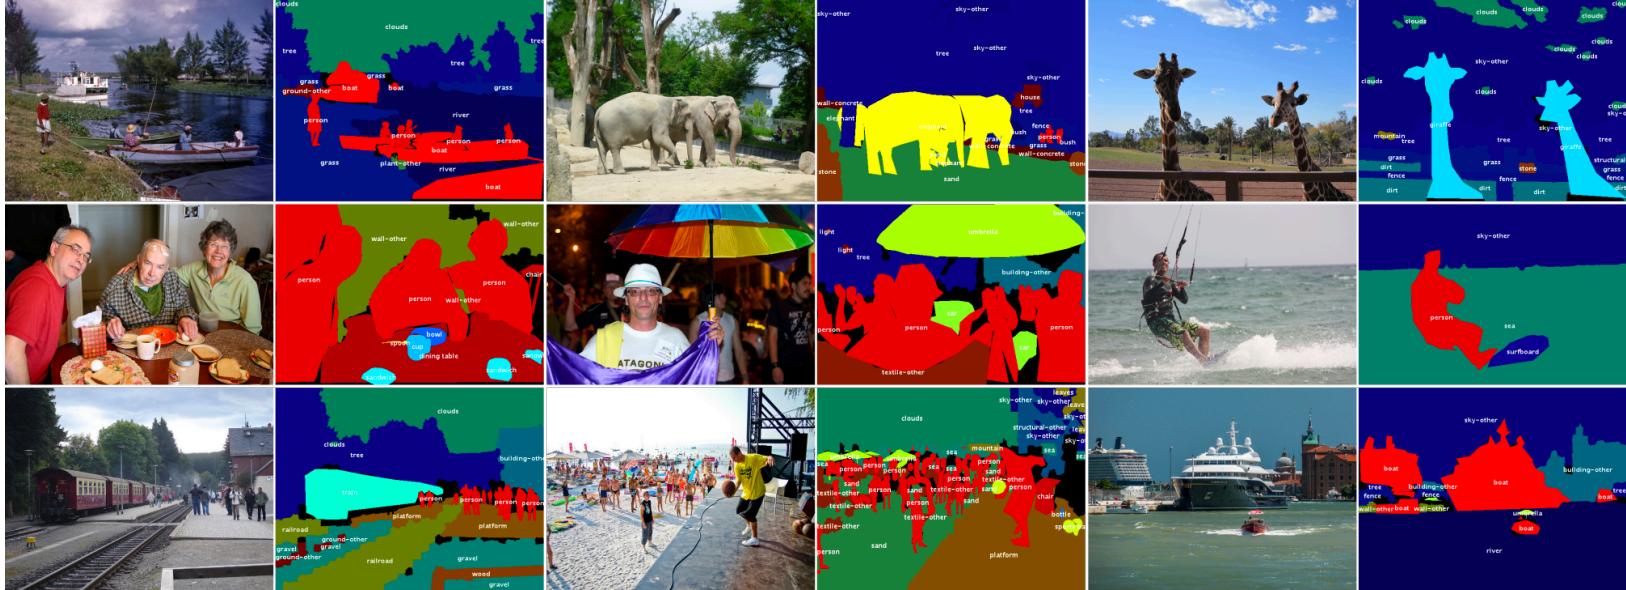


# MNIST RNN Classifier

```
class LSTM_Classifier(torch.nn.Module):
    def __init__(self, classes=10, input_size=28, hidden_size=128):
        super(LSTM_Classifier, self).__init__()
        #####
        # TODO: Build a LSTM classifier
        #####
        self.LSTM = nn.LSTM(input_size=input_size, hidden_size=hidden_size)
        self.predictor = nn.Linear(hidden_size, classes)
        self.softmax = nn.Softmax(dim=2)
        #####
        for m in self.RNN.modules():
            if isinstance(m, nn.Linear):
                nn.init.normal_(m.weight, mean=0, std=0.1)
                nn.init.constant_(m.bias, val=0)
        #####
        def forward(self, x):
            _, (h, c) = self.LSTM(x)
            y = self.predictor(h)
            prediction = self.softmax(y)

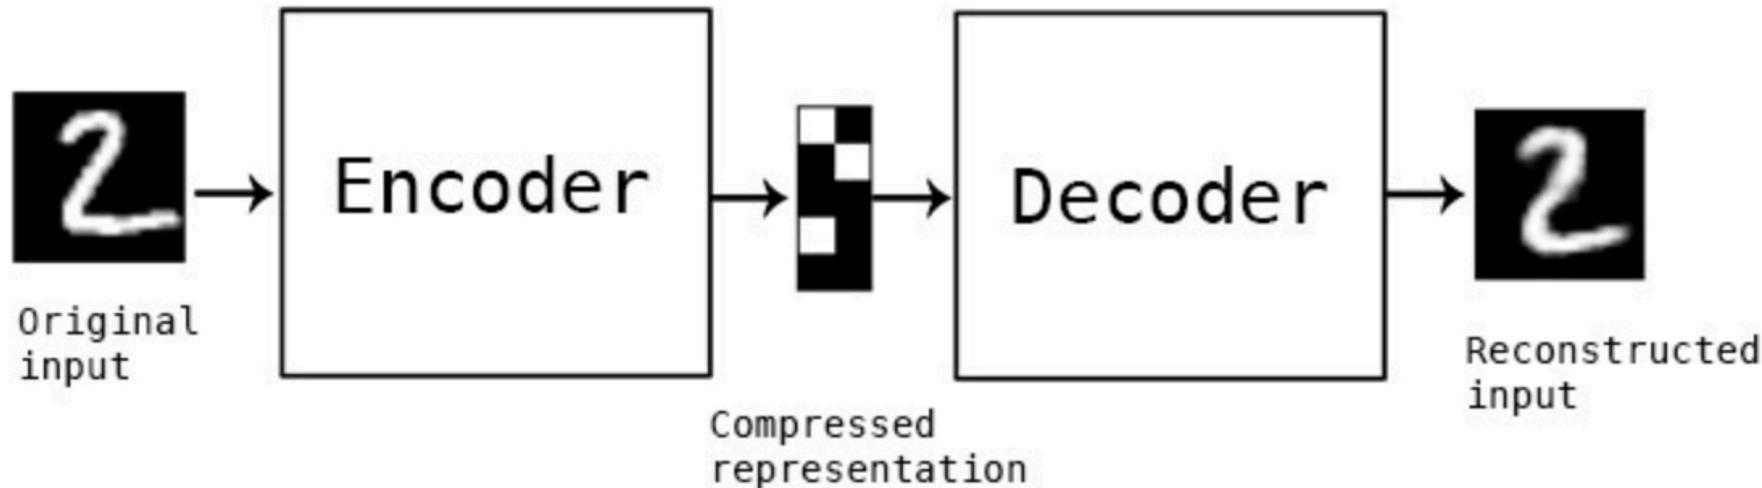
            return prediction.squeeze(0)
```

# Segmentation



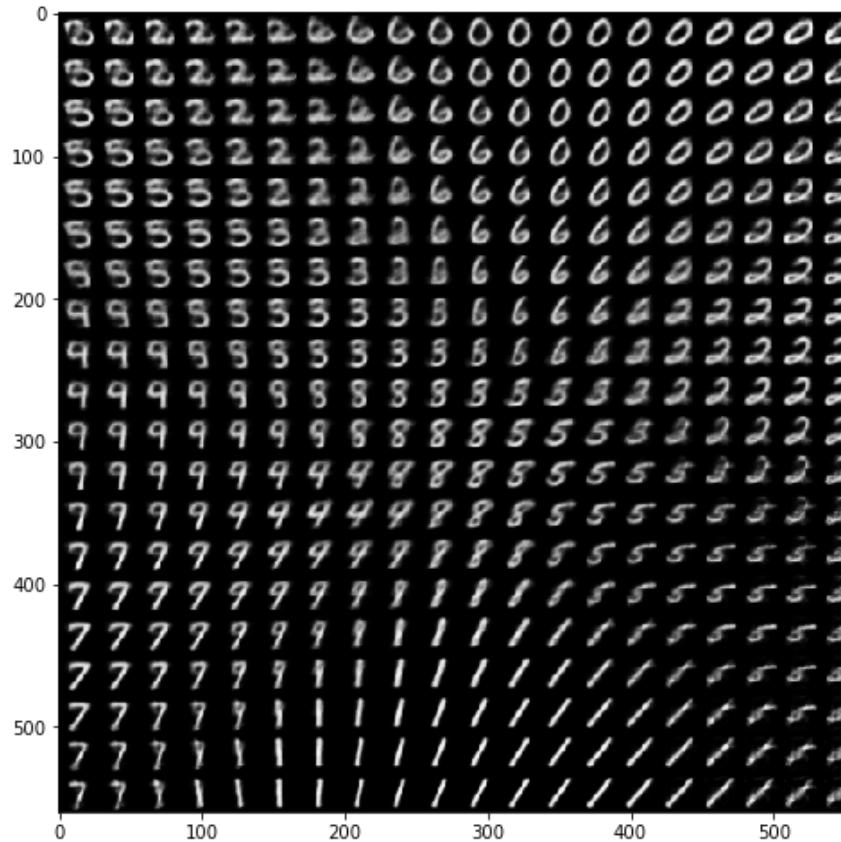
- Pixel-wise segmentation
- $Seg : \mathbb{R}^{n \times m \times 3} \rightarrow \mathbb{R}^{n \times m \times C}$

# Autoencoder architecture

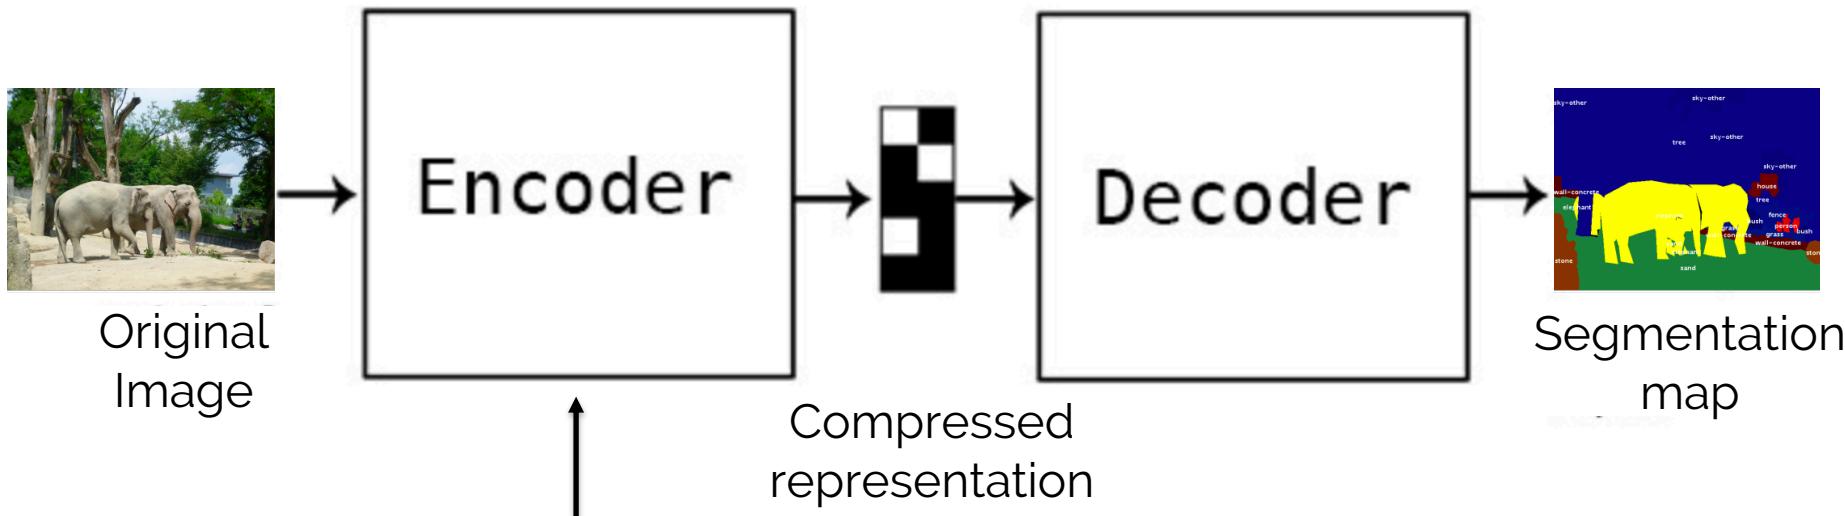




# Latent Space Interpolation (Variational Autoencoders)



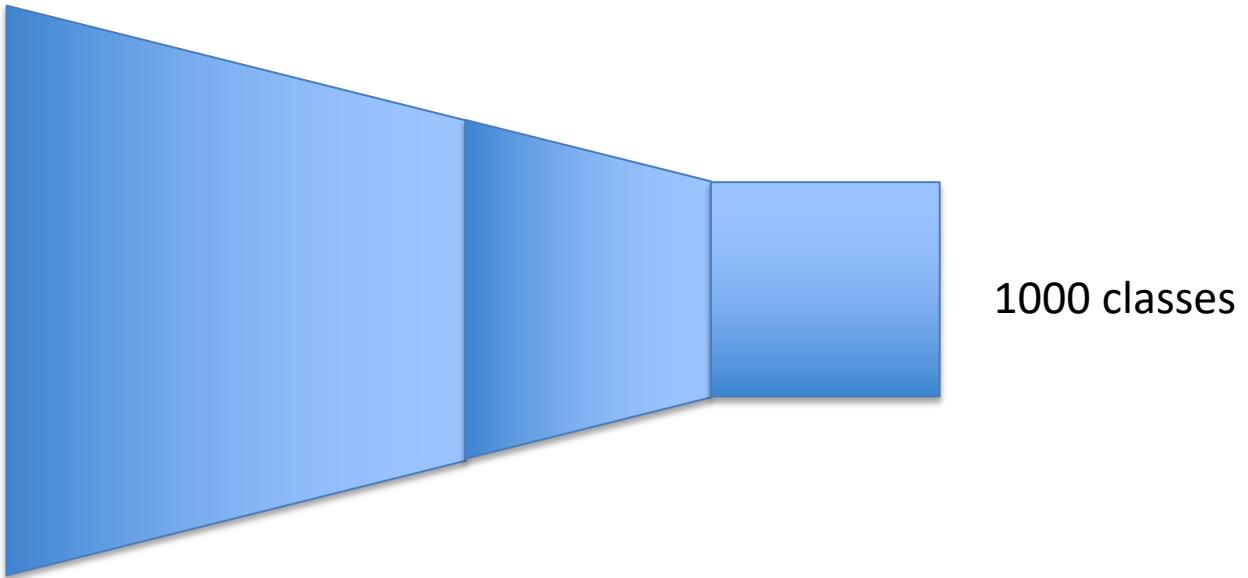
# Vanilla Segmentation architecture



- Encoder extracts features of images
  - Image features can be independent of task

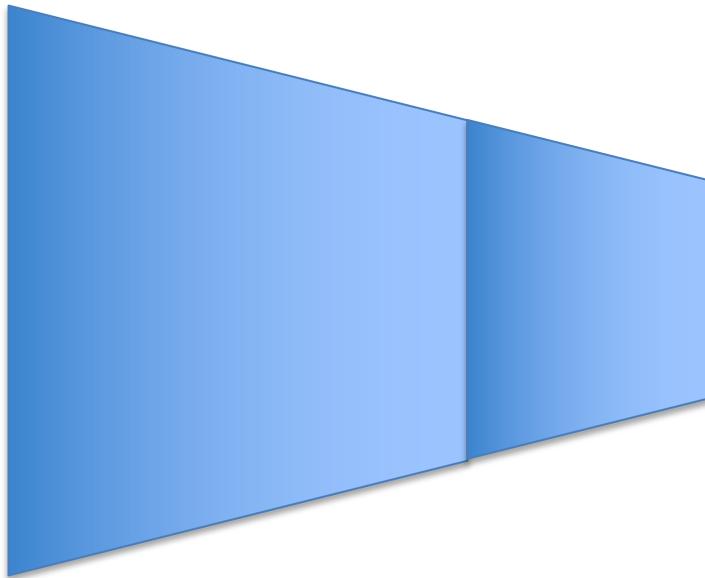
# Idea: Transfer Learning

# Transfer Learning



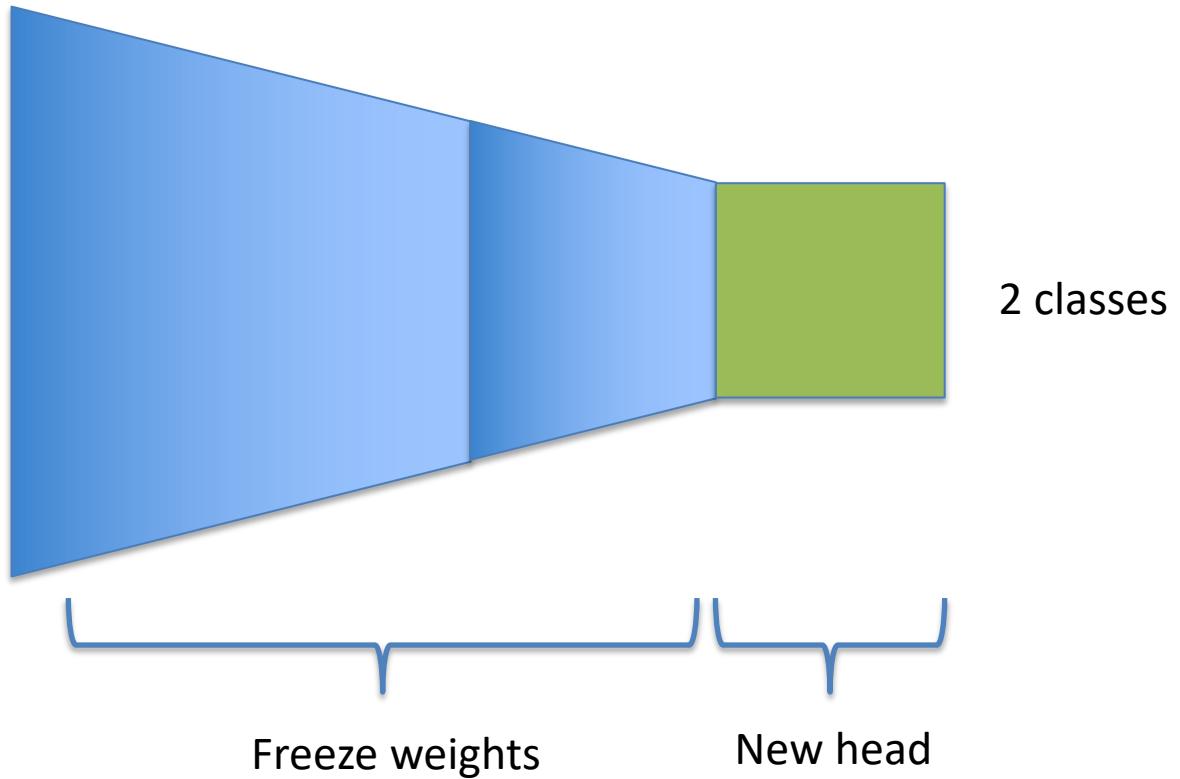
1000 classes

# Transfer Learning

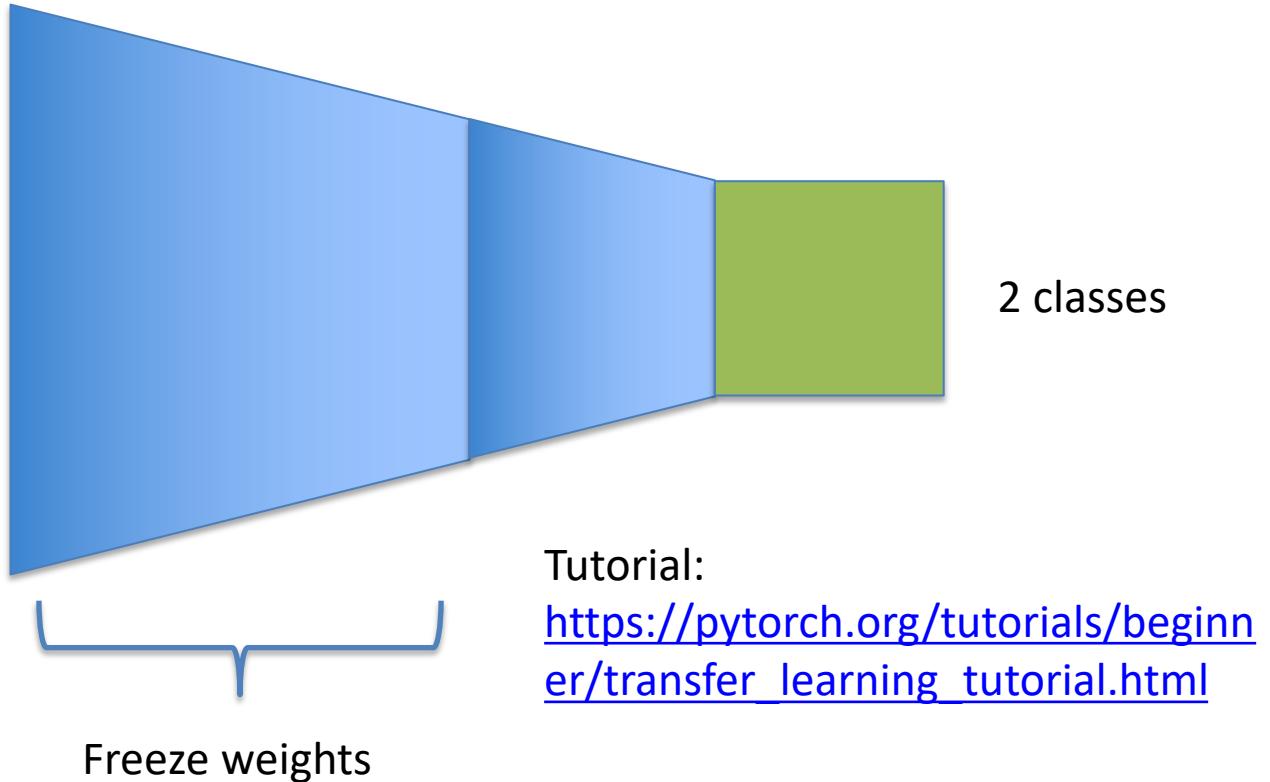


2 classes

# Transfer Learning: (Pre-)train



# Transfer Learning: Train



# Transfer Learning

- Idea:
  - take network trained on a similar task **with weights**
  - Add/Remove layers to make it suitable for your task
  - Two options:
    - Define different learning rates for pretrained part and new parts
    - Freeze pretrained part for some epochs (recommended)
- Advantage:
  - Don't have to learn early convolutional weights who don't really care about the exact image contents
  - Able to keep filters in the latter part of the network for new task

# Segmentation Exercise

```
class SegmentationNN(nn.Module):

    def __init__(self, num_classes=23):
        super(SegmentationNN, self).__init__()

        #####
        #                                     YOUR CODE
        #####
        from torchvision import models
        self.features = models.alexnet(pretrained=True).features
        self.classifier = nn.Sequential(
            nn.Dropout(),
            nn.Conv2d(256, 4096, kernel_size=1, padding=0),
            nn.ReLU(inplace=True),
            nn.Dropout(),
            nn.Conv2d(4096, 4096, kernel_size=1, padding=0),
            nn.ReLU(inplace=True),
            nn.Conv2d(4096, num_classes, kernel_size=1, padding=0),
            nn.Upsample(scale_factor=40),
            nn.Conv2d(num_classes, num_classes, kernel_size=3, padding=1),
        )
```

# Segmentation Exercise

How to do it:

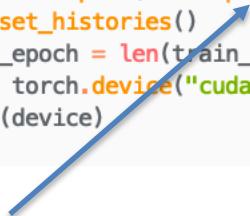
- Use a pretrained CNN (i.e. AlexNet) as feature extractor
- Add some convolutional layers to train decoder
- Upscale to output dimension
  - Here: simple upsampling

```
solver = Solver(optim_args={"lr": 1e-4, "weight_decay": 0.001},  
                 loss_func=torch.nn.CrossEntropyLoss(size_average=True, ignore_index=-1))  
solver.train(model, train_loader, val_loader, log_nth=5, num_epochs=5)
```

# How to freeze weights?

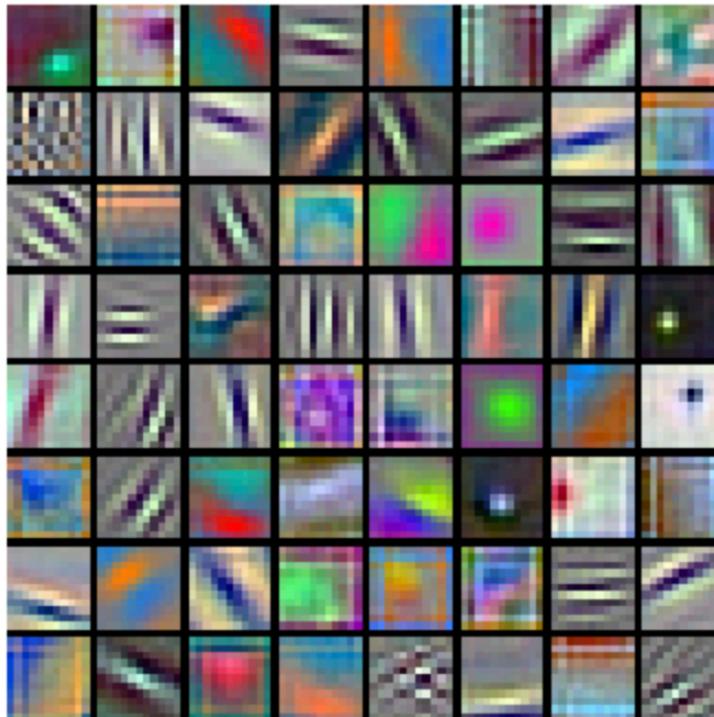
```
def train(self, model, train_loader, val_loader, num_epochs=10, log_nth=0):
    """
    Train a given model with the provided data.

    Inputs:
    - model: model object initialized from a torch.nn.Module
    - train_loader: train data in torch.utils.data.DataLoader
    - val_loader: val data in torch.utils.data.DataLoader
    - num_epochs: total number of training epochs
    - log_nth: log training accuracy and loss every nth iteration
    """
    optim = self.optim(model.parameters(), **self.optim_args)
    self._reset_histories()
    iter_per_epoch = len(train_loader)
    device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
    model.to(device)
```

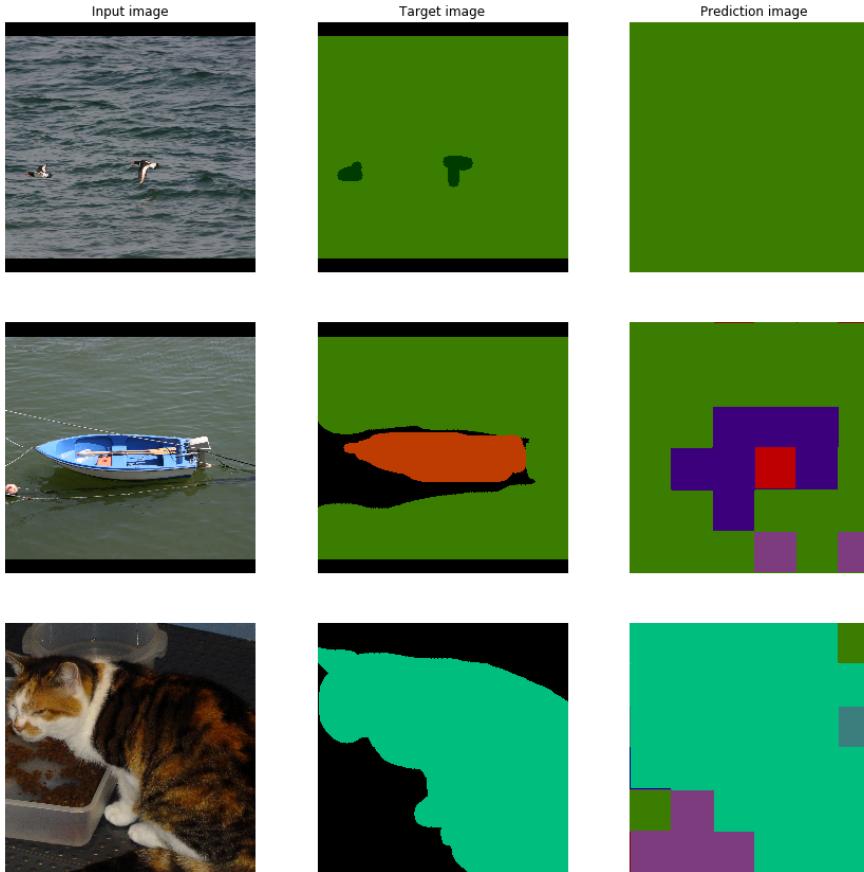


model.classifier.parameters()

# Visualize filter weights



# Outputs...



# Best performing student model

---

## Exercise Submission for I2DL

| Submission 0        | Submission 1 | Submission 2        | Submission 3 | Submission 4 | Submission 5 | Submission 6 | Submission 7 | Bonus |
|---------------------|--------------|---------------------|--------------|--------------|--------------|--------------|--------------|-------|
| Rank                | User         | Date submitted      |              |              | Score        | Pass         |              |       |
| <b>Submission 5</b> |              |                     |              |              |              |              |              |       |
| #1                  | w0256        | 2020-01-14 07:42:59 |              |              | 93.37        | ✓            |              |       |
| #2                  | w0220        | 2020-01-22 01:43:11 |              |              | 92.61        | ✓            |              |       |
| #3                  | w0445        | 2019-12-26 16:17:56 |              |              | 92.49        | ✓            |              |       |

# What do they do better?

Chats #####

# YOUR CODE #####

from torchvision.models.segmentation.deeplabv3 import DeepLabHead  
from torchvision import datasets, models, transforms  
device = torch.device("cuda:0" if torch.cuda.is\_available() else "cpu")

self.deeplab = models.segmentation.deeplabv3\_resnet101(pretrained=True, progress=True)  
# Change output size  
self.deeplab.classifier = DeepLabHead(2048, num\_classes)  
self.deeplab.to(device)

Rudi, Simply Simon, pyr-0-... 00:01  
https://twitter.com/ndnink120/s/

# END OF YOUR CODE #####



# Practical (Pytorch) Tips

# Coding Project Outline

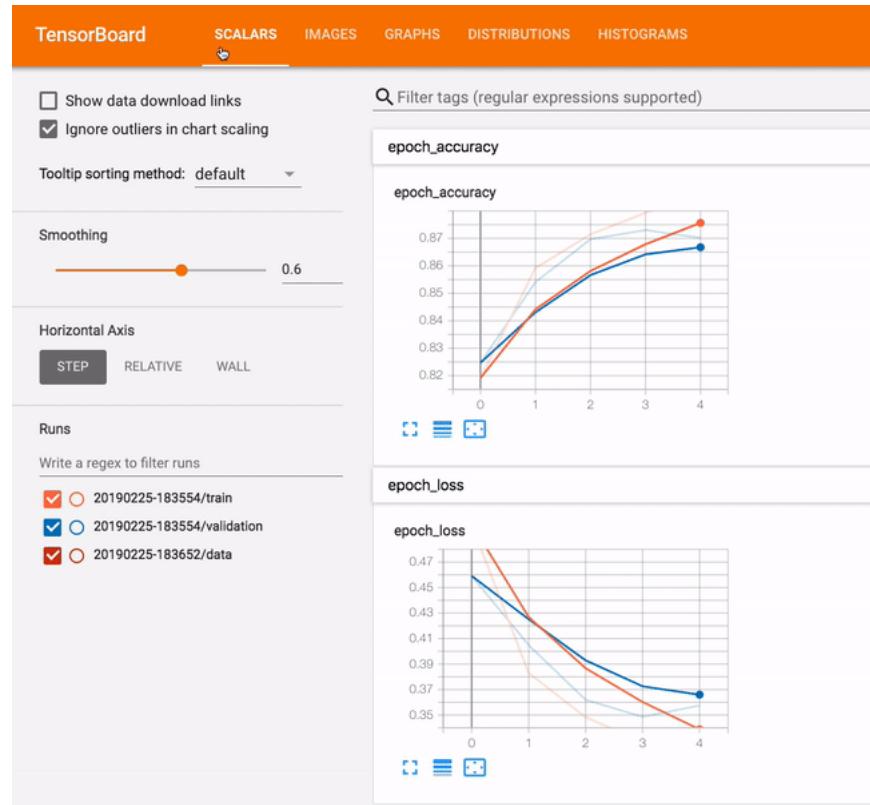
- Data
  - Write visualizer
  - Create data loader and unit test it
- Solver (main program)
  - Adopt from previous project or use existing libraries
  - Log all input variables
  - Always display all relevant graphs and numbers throughout the training process
    - Use existing tools like tensorboard

# Coding Project Outline

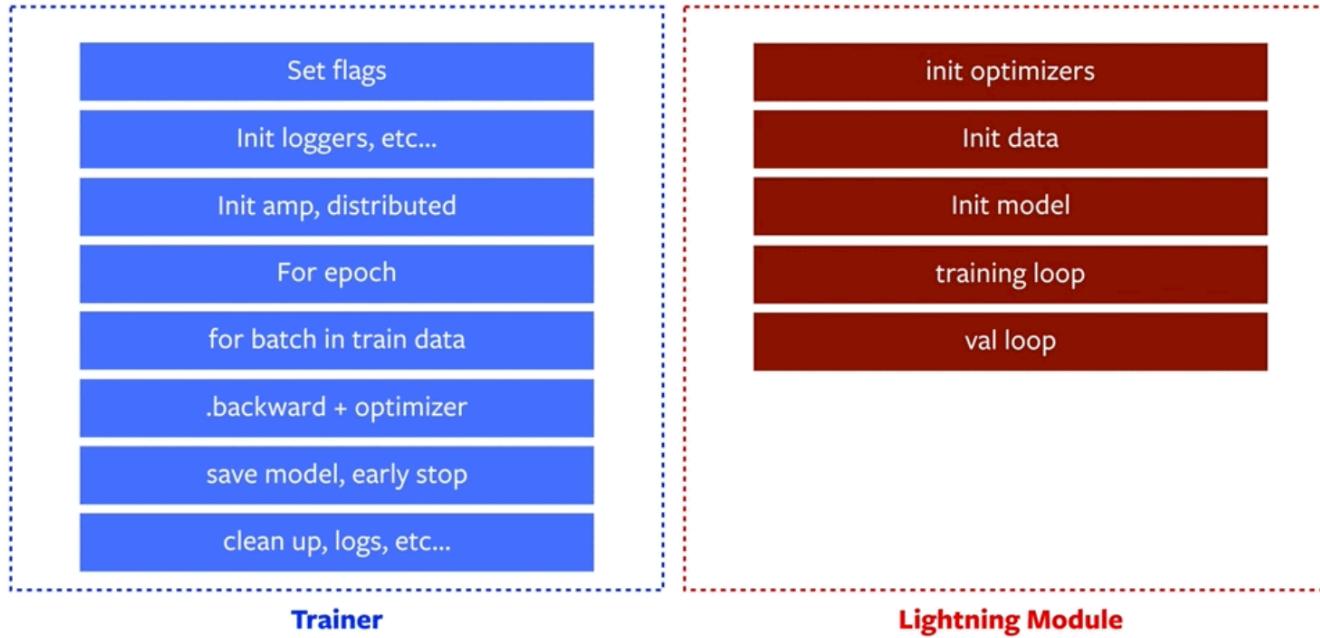
- Network
  - Start **very simple** (no pretrained model)
  - Overfit to a low number of examples to check dataloader (solver) and network
  - Expand from there and always test multiple hyperparameters
- Other stuff
  - Implement data augmentation (transforms)

# How to visualize training?

- Use libraries like
  - Visdom
  - Tensorboard
    - Use pytorch wrappers like tensorboardX etc.



# Pre-existing solvers?



Link: <https://github.com/PyTorchLightning/pytorch-lightning>

# Exam

# Exam: Hard facts

- 90 minutes to achieve 90 points
- No calculator, laptop or script allowed
- 3 parts:
  - Multiple choice questions
  - Short questions
  - 4 long questions
- Covers lecture and exercise content
- No explicit coding but pseudo code possible

# Notes

- TUM Exam
  - There will be no official classroom review but everything will be done online
  - We will outline the necessary steps after the exam on moodle
- Read the instructions on page 1 carefully!

# Notes

- Multiple Choice
  - Theme can change but will probably be as outlined in the mock exam
- Mock Exam
  - Solutions will be posted today
  - There can be errors in the solutions
  - all discussions on moodle

# Topics you should know!

List not exhaustive!

- **Lecture 1: Intro**
  - Motivation and Applications
- **Lecture 2: Machine Learning Basics**
  - (Un-)Supervised Learning
  - (K-)Nearest Neighbors
  - Data Splitting + Cross Validation
  - (Linear/Logistic) Regression
- **Lecture 3: Intro to NNs**
  - Activation Functions
  - Loss Functions
  - Computational Graphs
- **Lecture 4: Backpropagation**
  - NNs as Computational Graphs
  - Backpropagation
  - Gradient Descent
- **Lecture 5: Optimization for NNs**
  - Stochastic Gradient Descent and deviations
  - Convergence of GD/ SGD
  - Newton's and other 2nd order methods
- **Lecture 6: Training NNs I**
  - Over- and Underfitting
  - Interpretation of train/ val curves
- **Lecture 7: Training NNs II**
  - Activation Functions and Loss Functions
  - Weight Initialization
- **Lecture 8: Advanced Regularization techniques**
  - Data Augmentation
  - Regularization (Normalization, Dropout, ...)
- **Lecture 9: Intro to CNNs**
  - Convolutions (Forward passes, Dimensions)
  - Pooling
- **Lecture 10: CNN Architectures**
  - Receptive Field
  - Different Architectures (LeNet, AlexNet, VGGNet, ResNet, Inception Layer, EfficientNet, R-CNNs, FCNs, U-Net)
- **Lecture 11: RNNs**
  - Transfer Learning
  - RNNs & LSTMs
    - Dependencies and Vanishing Gradient
    - Shortcomings of RNNs
    - Improvements via LSTM cells
- **Lecture 12: Autoencoder and GANs**
  - Encoder – Decoder Architecture with Use-Cases
  - Variational Autoencoder
  - GANs with applications

# Organization

# Organization

- Office hours
  - Will continue until the exam
- Next lecture:
  - Guest lecture with Timo Aila



Questions?

Good luck for the exam

19.02.20 – 13:30h